

Advanced Static Analysis of Atomicity in Concurrent Programs through Facebook Infer

Excel@FIT 2021 — Paper No. 57

Dominik Harmim

Supervisor: prof. Ing. Tomáš Vojnar, Ph.D.

xharmi00@stud.fit.vutbr.cz

Brno University of Technology, Faculty of Information Technology



6th May 2021

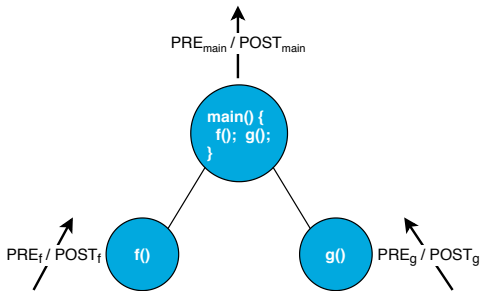
- Detecting and checking desired **atomicity** of **function call sequences**.
 - Often required in **concurrent programs**.
 - **Violation** may cause **nasty errors**.

```
void invoke(char *method) {  
    ...  
    if (server.is_registered(method)) {  
        server.invoke(method);  
    }  
    ...  
}
```

The sequence of **is_registered** and **invoke** should be **executed atomically**.

If **not locked**, the method can be unregistered by a **concurrent thread**.

- Open-source **static analysis framework** for **interprocedural analyses**.
 - Based on **abstract interpretation**.
- Highly **scalable**.
 - Follows principles of **compositionality**.
 - Computes function **summaries** **bottom-up** on call-trees.
- Supports C, C++, Java, Obj-C, C#.

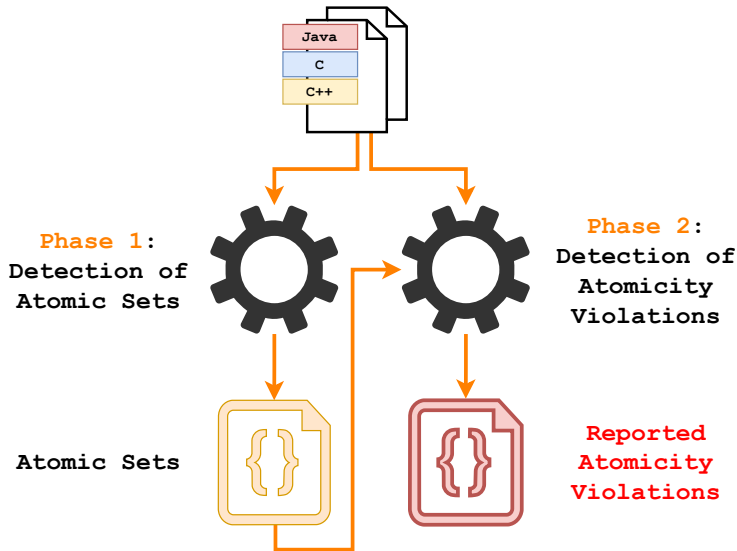


- Facebook Infer plugin created within the author's BSc thesis:



HARMIM, D. *Static Analysis Using Facebook Infer to Find Atomicity Violations*. Brno, 2019. Bachelor's thesis. Brno University of Technology, Faculty of Information Technology. Supervisor VOJNAR, T.

- **Assumption:** call sequences executed atomically once should (probably) be executed always atomically.
- Implemented for C programs that use PThread locks.
- Limited scalability on extensive codebases.
- Reports many false alarms when analysing real-life code.



1 Detection of **atomic call sets**.

- Approximates sequences by sets.
- Summary:** $\chi \in 2^\Sigma \times 2^{2^\Sigma}$
(set of all calls, set of atomic call sets)

```
void f() {
    lock(L);
    x(); y(); z(); // x.y.z -> {x,y,z}
    unlock(L);
    a();
    lock(L);
    z(); y(); x(); // z.y.x -> {x,y,z}
    unlock(L);
}
```

$$\chi_f = (\{a, x, y, z\}, \{\{x, y, z\}\})$$

$$\chi'_f = (x \cdot y \cdot z \cdot a, \{x \cdot y \cdot z, z \cdot y \cdot x\})$$

2 Detection of **atomicity violations**.

- Derives "**atomic pairs**" from the first phase: $\Omega \in 2^{\Sigma \times \Sigma}$
- Looks for **non-atomic pairs** of calls assumed to **run atomically**.
- Summary:** $\chi \in 2^{\Sigma \times \Sigma}$
(set of atomicity violations)

```
void g() {
    a(); x(); y(); b();
}
```

$$\Omega = \{(x, y), (x, z), (y, x), (y, z), (z, x), (z, y)\}$$

$$\Omega' = \{(x, y), (y, z), (z, y), (y, x)\}$$

$$(x, y) \in \Omega \implies \chi_g = \{(x, y)\}$$

- Support for C++ and Java.
 - Working with advanced locks: re-entrant locks, monitors, lock guards, etc.
- Distinguishing different lock instances.
 - Approximating lock objects using syntactic access paths — a representation of heap locations via the paths used to access them.
- Analysis's parametrisation:
 - ignoring generic functions — concentrating on critical functions;
 - limiting the number of calls or the depth of nested calls in critical sections.

- **Scalability** evaluated on 54 real-life complex C programs.
 - 806,431 LOC in total.
- **Double acceleration** in average.

	v1.0.0		v2.0.0	
	Phs. 1	Phs. 2	Phs. 1	Phs. 2
Avg. Time (s)	70.98	109.11	37.96	50.93
Total Time (s)	4,117	5,892	2,164	2,750


- Experiments with **Apache Cassandra** and **Apache Tomcat** (both ~ 250 KLOC).
 - Successfully **rediscovered** already fixed reported real bugs.
 - The number of reported bugs was **significantly reduced** ($\sim 4\times$).
 - Still hard to say which of the bugs are real — the **accuracy** needs to be further improved.


Atomer's Extensions:

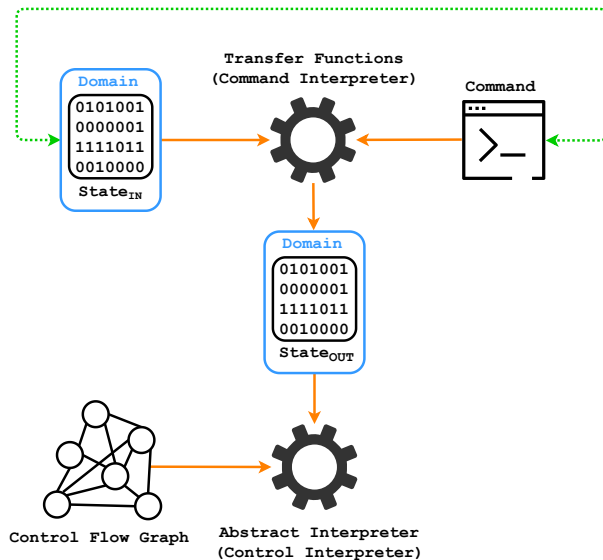
- Proposed and implemented extensions for Atomer:
 - approximation with sets, support for C++ and Java, distinguishing different lock instances, parametrisation of the analysis.
- Successfully tested and experimentally evaluated.
 - Both scalability and accuracy were increased.
- Experiments with real-life programs.

Future goals:

- Increase accuracy/reduce the number of false alarms:
 - Combining with dynamic analysis.
 - Statistic ranking of atomic functions/reported errors.
 - Considering formal parameters of function.
 - Machine learning of analysis parameter values.







Real-life bug in package `org.apache.catalina.core.StandardContext`.

```
public void addParameter(String name, String value) {  
    ...  
    if (parameters.get(name) != null)  
        throw new IllegalArgumentException  
            (sm.getString("standardContext.parameter.duplicate", name));  
  
    // Add this parameter to our defined set  
    synchronized (parameters) {  
        parameters.put(name, value);  
    }  
    fireContainerEvent("addParameter", name);  
}
```

- Access path used for locks identification: $\pi \in \Pi ::= Var \times Field^*$
 - Var is a set of all variables,
 - $Field$ is a set of field names.
- Identification of critical sections: $(\pi, l) \in \Pi \times \mathbb{N}^T$
 - π is an access path that identifies the lock object that locks the section,
 - l is the number of locks of the lock object identified by π ,
 - \mathbb{N}^T denotes $\mathbb{N} \cup \{T\}$
 - T represents a number larger than some upper bound $t \in \mathbb{N}$.
- Representation of lock guards: $(\pi_g, L) \in \Pi \times 2^\Pi$
 - π_g is an access path that identifies the lock guard,
 - L is a set of access paths that identify the lock objects associated with the guard identified by π_g .